

# Statistical Learning/ Lab 4

Sascha Strobl

5/16/2018

## Support Vector Machine classification and hierarchical clustering

Load dataset and remove two columns:

```
setwd("~/Documents/CGUClasses/ZOLDClasses/Statistical Learning/Lab4");
GH = read.table("glendalehousing.txt", header = T);
drops <- c("Address", "LotArea.1")
GH1=GH[ , !(names(GH) %in% drops)]
```

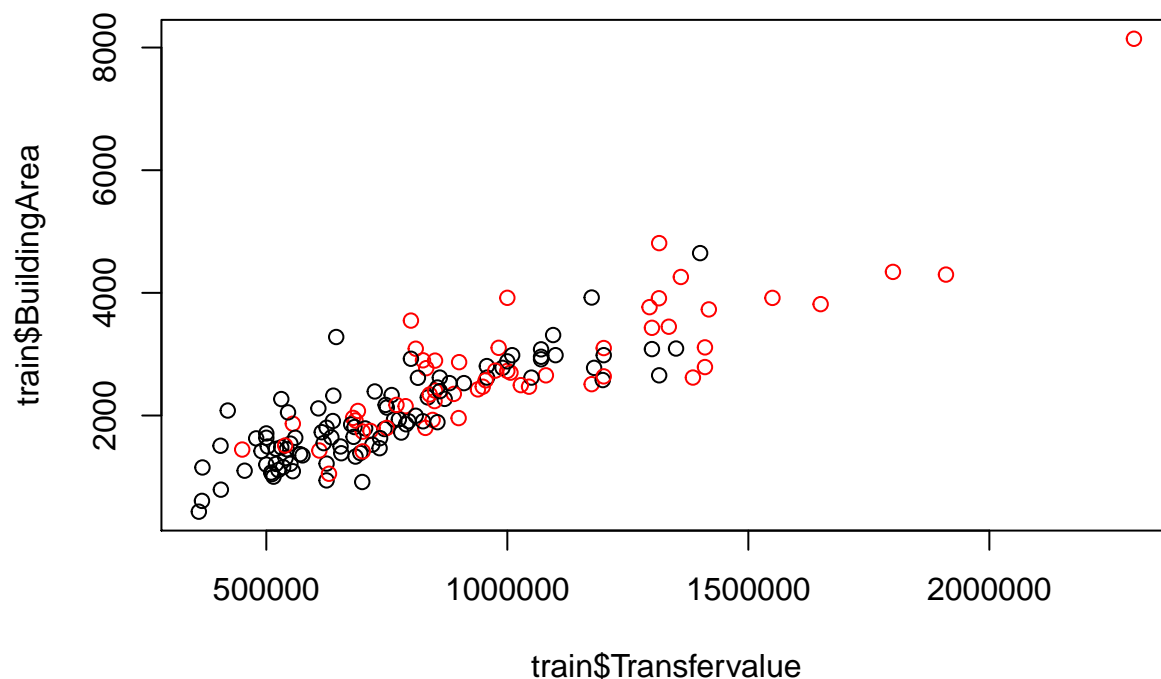
Get dimensions of dataset:

```
dim(GH1)
```

```
## [1] 173 12
```

Run Support Vector Machine classification with three different kernels (linear, polynomial, radial):

```
keeps <- c("Transfervalue", "BuildingArea", "Pool")
GH2=GH1[keeps]
train=GH2[1:150,]
test=GH2[151:173,]
train$Pool[train$Pool == 0] <- -1
test$Pool[test$Pool == 0] <- -1
train$Pool=as.factor(train$Pool);
test$Pool=as.factor(test$Pool);
plot(x=train$Transfervalue, y=train$BuildingArea, col=factor(train$Pool));
```



```

library(e1071);
#Linear
set.seed(1);
tune.out=tune(svm,Pool~.,data=train,kernel="linear",ranges=list(cost=c(0.001,0.01,0.1,1,5,10,100)));
summary(tune.out);

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##     1
##
## - best performance: 0.3333333
##
## - Detailed performance results:
##   cost      error dispersion
## 1 1e-03 0.3733333 0.11417985
## 2 1e-02 0.3533333 0.09962894
## 3 1e-01 0.3400000 0.10159226
## 4 1e+00 0.3333333 0.09938080
## 5 5e+00 0.3333333 0.09938080
## 6 1e+01 0.3333333 0.09938080
## 7 1e+02 0.3333333 0.09938080

bestmod =tune.out$best.model;
summary(bestmod);

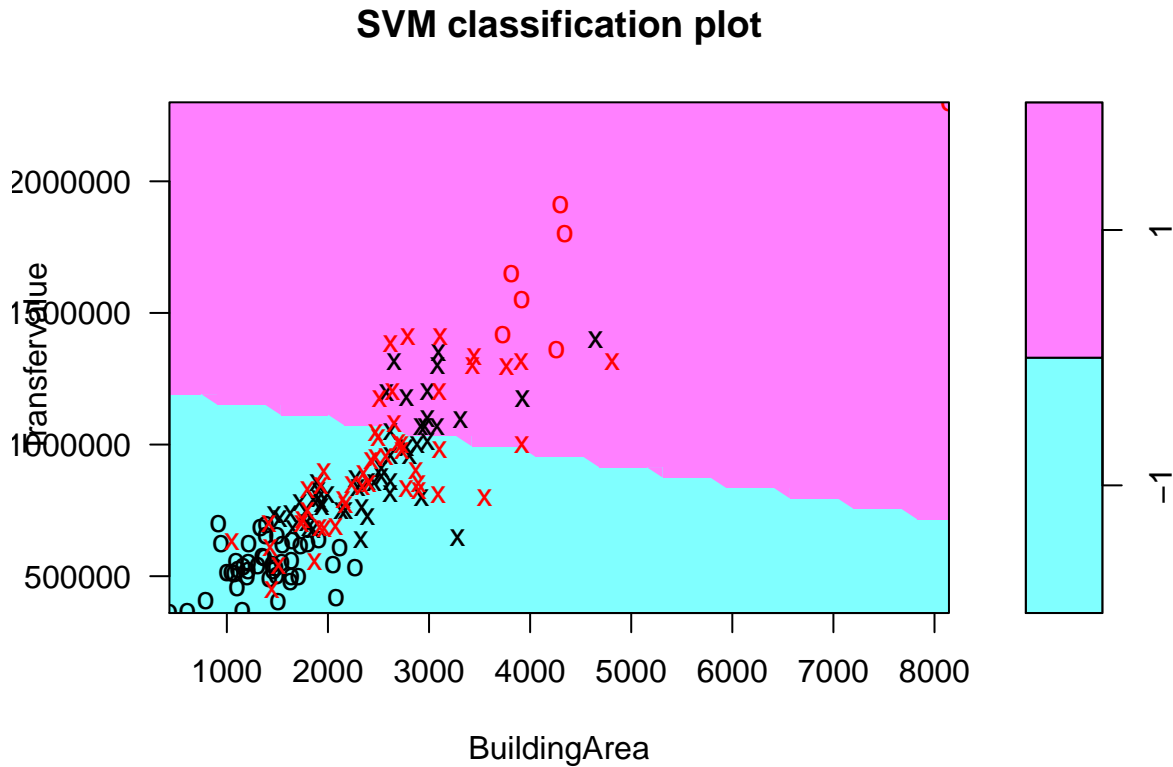
##
## Call:
## best.tune(method = svm, train.x = Pool ~ ., data = train, ranges = list(cost = c(0.001,
##   0.01, 0.1, 1, 5, 10, 100)), kernel = "linear")
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel:  linear
##         cost:  1
##        gamma:  0.5
##
## Number of Support Vectors:  99
##
##   ( 50 49 )
##
##
## Number of Classes:  2
##
## Levels:
##   -1 1

ypred=predict(bestmod ,test);
table(predict =ypred , truth= test$Pool);

```

```
##      truth
## predict -1  1
##      -1 13  6
##      1  2  2
```

```
plot(bestmod,train)
```



```
#Polynomial
set.seed(1);
tune.out=tune(svm,Pool~.,data=train,kernel="polynomial",ranges=list(cost=c(0.001,0.01,0.1,1,5,10,100)))
summary(tune.out);
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##   0.1
##
## - best performance: 0.2933333
##
## - Detailed performance results:
##   cost      error dispersion
## 1 1e-03 0.3666667 0.10540926
## 2 1e-02 0.3333333 0.10423146
## 3 1e-01 0.2933333 0.11417985
## 4 1e+00 0.3066667 0.09532271
## 5 5e+00 0.3133333 0.09962894
## 6 1e+01 0.3200000 0.11243654
```

```

## 7 1e+02 0.3200000 0.11243654
bestmod =tune.out$best.model;
summary(bestmod);

##
## Call:
## best.tune(method = svm, train.x = Pool ~ ., data = train, ranges = list(cost = c(0.001,
##      0.01, 0.1, 1, 5, 10, 100)), kernel = "polynomial")
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: polynomial
##      cost:  0.1
##   degree:  3
##   gamma:  0.5
##   coef.0:  0
##
## Number of Support Vectors:  100
##
##   ( 50 50 )
##
##
## Number of Classes:  2
##
## Levels:
##   -1 1

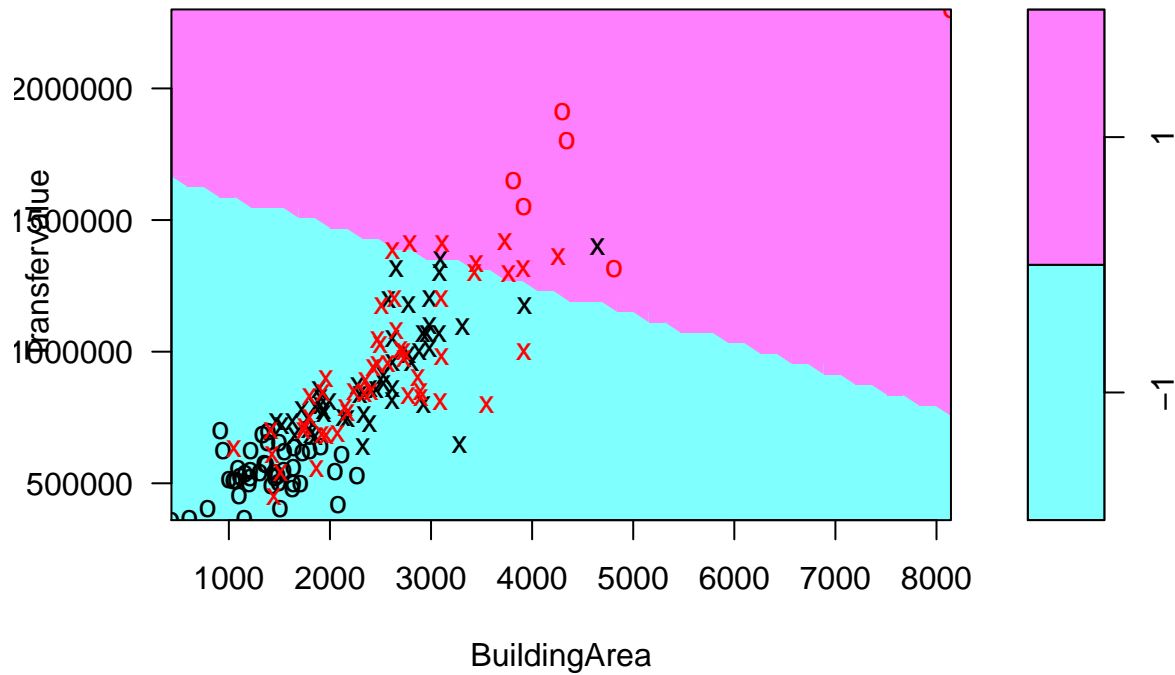
ypred=predict (bestmod ,test);
table(predict =ypred , truth= test$Pool);

##      truth
## predict -1  1
##      -1 15  6
##      1  0  2

plot(bestmod,train)

```

## SVM classification plot



```
#Radial
set.seed(1);
tune.out=tune(svm,Pool~.,data=train,kernel="radial",ranges
              =list(cost=c(0.001,0.01,0.1,1,5,10,100)));
summary(tune.out);
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##     1
##
## - best performance: 0.3466667
##
## - Detailed performance results:
##   cost      error dispersion
## 1 1e-03 0.3733333 0.1141798
## 2 1e-02 0.3733333 0.1141798
## 3 1e-01 0.3800000 0.1090928
## 4 1e+00 0.3466667 0.1079552
## 5 5e+00 0.3533333 0.1177987
## 6 1e+01 0.3733333 0.1003697
## 7 1e+02 0.3533333 0.1297671
```

```
bestmod =tune.out$best.model;
summary(bestmod);
```

```
##
```

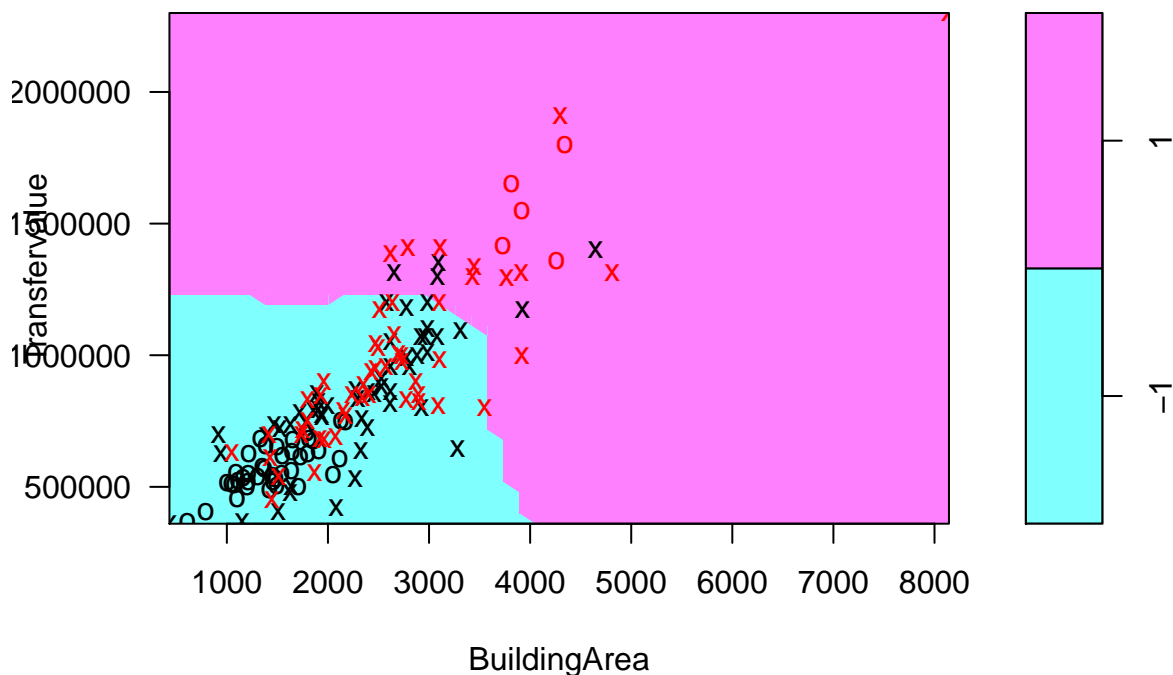
```
## Call:
## best.tune(method = svm, train.x = Pool ~ ., data = train, ranges = list(cost = c(0.001,
## 0.01, 0.1, 1, 5, 10, 100)), kernel = "radial")
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: radial
##     cost:    1
##    gamma:    0.5
##
## Number of Support Vectors: 105
##
## ( 54 51 )
##
##
## Number of Classes: 2
##
## Levels:
## -1 1

ypred=predict (bestmod ,test);
table(predict =ypred , truth= test$Pool);

##      truth
## predict -1  1
##      -1 15  6
##       1  0  2


plot(bestmod,train)
```

**SVM classification plot**



Hierarchical Clustering on the USArrests dataset:

```
data(USArrests);  
clusters <- hclust(dist(USArrests))  
plot(clusters)
```



```
x=USArrests;  
hc.complete=hclust(dist(x),method ="complete");  
hc.average=hclust(dist(x),method ="average");  
hc.single=hclust(dist(x),method ="single");  
par(mfrow =c(1,3));  
plot(hc.complete,main="Complete;Linkage",xlab="",sub="",cex=.9);  
plot(hc.average,main ="Average Linkage",xlab="",sub="",cex=.9);  
plot(hc.single,main ="Single Linkage",xlab="",sub="",cex=.9);
```

